


In statistics, [Kolmogorov–Smirnov test](#) is a popular procedure to test, from a sample $\{x_1, \dots, x_n\}$ is drawn from a distribution (F) , or usually (F_{θ_0}) , where (F_{θ}) is some parametric distribution. For instance, we can test $(H_0: X_i \sim \mathcal{N}(0,1))$ (where $(\theta_0 = (\mu_0, \sigma_0^2) = (0,1))$) using that test. More specifically, I wanted to discuss today (p) -values. Given (n) let us draw $(\mathcal{N}(0,1))$ samples of size (n) , and compute the (p) -values of [Kolmogorov–Smirnov tests](#)

```
n=300
p = rep(NA,1e5)
for(s in 1:1e5){
  X = rnorm(n,0,1)
  p[s] = ks.test(X,"pnorm",0,1)$p.value
}
```

We can visualise the distribution of the (p) -values below (I added some Beta distribution fit here)

```
library(fitdistrplus)
fit.dist = fitdist(p,"beta")
hist(p,probability = TRUE,main="",xlab="",ylab="")
vu = seq(0,1,by=.01)
vv = dbeta(vu,shape1 = fit.dist$estimate[1], shape2 =
fit.dist$estimate[2])
lines(vu,vv,col="dark red", lwd=2)
```



It looks like it is quite uniform (theoretically, the (p) -value is uniform). More specifically, the (p) -value was lower than 5% in 5% of the samples

[note: here I compute 'mean(p<=.05)' but I have some trouble with the '<' and '>' symbols, as always]

```
mean(p<=.05)
[1] 0.0479
```


i.e. we wrongly reject $(H_0: X_i \sim \mathcal{N}(0,1))$ is 5% of the samples.

As discussed previously on the blog, in many cases, we do care about the distribution, and not really the parameters, so we wish to test something like $(H_0: X_i \sim \mathcal{N}(\mu, \sigma^2))$, for some (μ) and (σ^2) . Therefore, a natural idea can be to test $(H_0: X_i \sim \mathcal{N}(\hat{\mu}, \hat{\sigma}^2))$, for some estimates of (μ) and (σ^2) . That's the idea of [Lilliefors test](#). More specifically, [Lilliefors test](#) suggests to use [Kolmogorov–Smirnov](#) statistics, but corrects the (p) -value. Indeed, if we draw many samples, and use [Kolmogorov–Smirnov](#) statistics and its classical (p) -value to test for $(H_0: X_i \sim \mathcal{N}(\hat{\mu}, \hat{\sigma}^2))$,

```
n=300
p = rep(NA,1e5)
for(s in 1:1e5){
  X = rnorm(n,0,1)
  p[s] = ks.test(X,"pnorm",mean(X),sd(X))$p.value
}
```

we see clearly that the distribution of (p) -values is no longer uniform

```
fit.dist = fitdist(p,"beta")
hist(p,probability = TRUE,main="",xlab="",ylab="")
vu = seq(0,1,by=.01)
vv = dbeta(vu,shape1 = fit.dist$estimate[1], shape2 =
fit.dist$estimate[2])
lines(vu,vv,col="dark red", lwd=2)
```



More specifically, if x_i 's are actually drawn from some Gaussian distribution, there are no chance to reject H_0 , the p -value being almost never below 5%

```
mean(p<=.05)
[1] 0.00012
```

Usually, to interpret that result, the heuristics is that $\hat{\mu}$ and $\hat{\sigma}^2$ are both based on the sample, while previously 0 and 1 were based on some prior knowledge. Somehow, it reminded me on the classical problem when mention when we introduce cross-validation, which is [Goodhart's law](#)


When a measure becomes a target, it ceases to be a good measure

i.e. we cannot assess goodness of fit using the same data as the ones used to estimate parameters. So here, why not use some [hold-out](#) (or cross-validation) procedure : split the dataset in two parts, $\{x_1, \dots, x_k\}$ (with k Kolmogorov–Smirnov statistics on it to test if x_i 's are drawn from some Gaussian distribution. More precisely, will the p -value computed using the standard [Kolmogorov–Smirnov](#) procedure be ok here. Here, I tried two scenarios, k/n being either $1/3$ or $2/3$),

```
p = matrix(NA, 1e5, 4)
for(s in 1:1e5){
  X = rnorm(n, 0, 1)
  p[s, 1] = ks.test(X, "pnorm", 0, 1)$p.value
  p[s, 2] = ks.test(X, "pnorm", mean(X), sd(X))$p.value
  p[s, 3] = ks.test(X[1:200], "pnorm", mean(X[201:300]), sd(X[201:300]))$p.value
  p[s, 4] = ks.test(X[201:300], "pnorm", mean(X[1:200]), sd(X[1:200]))$p.value
}
```


Again, we can visualize the distributions of p -values, in the case where $(1/3)$ of the data is used to estimate (μ) and (σ^2) , and $(2/3)$ of the data is used to test

```
fit.dist = fitdist(p[,3], "beta")
hist(p[,3], probability = TRUE, main="", xlab="", ylab="")
vu=seq(0,1,by=.01)
vv=dbeta(vu, shape1 = fit.dist$estimate[1], shape2 =
fit.dist$estimate[2])
lines(vu,vv,col="dark red", lwd=2)
```



and in the case where $(2/3)$ of the data is used to estimate (μ) and (σ^2) , and $(1/3)$ of the data is used to test

```
fit.dist = fitdist(p[,4], "beta")
hist(p[,4], probability = TRUE, main="", xlab="", ylab="")
vu=seq(0,1,by=.01)
vv=dbeta(vu, shape1 = fit.dist$estimate[1], shape2 =
fit.dist$estimate[2])
lines(vu,vv,col="dark red", lwd=2)
```



Observe here that we (wrongly) reject too frequently H_0 , since the p -values are below 5% in 25% of the scenarios, in the first case (less data used to estimate), and 9% of the scenarios, in the second case (less data used to test)

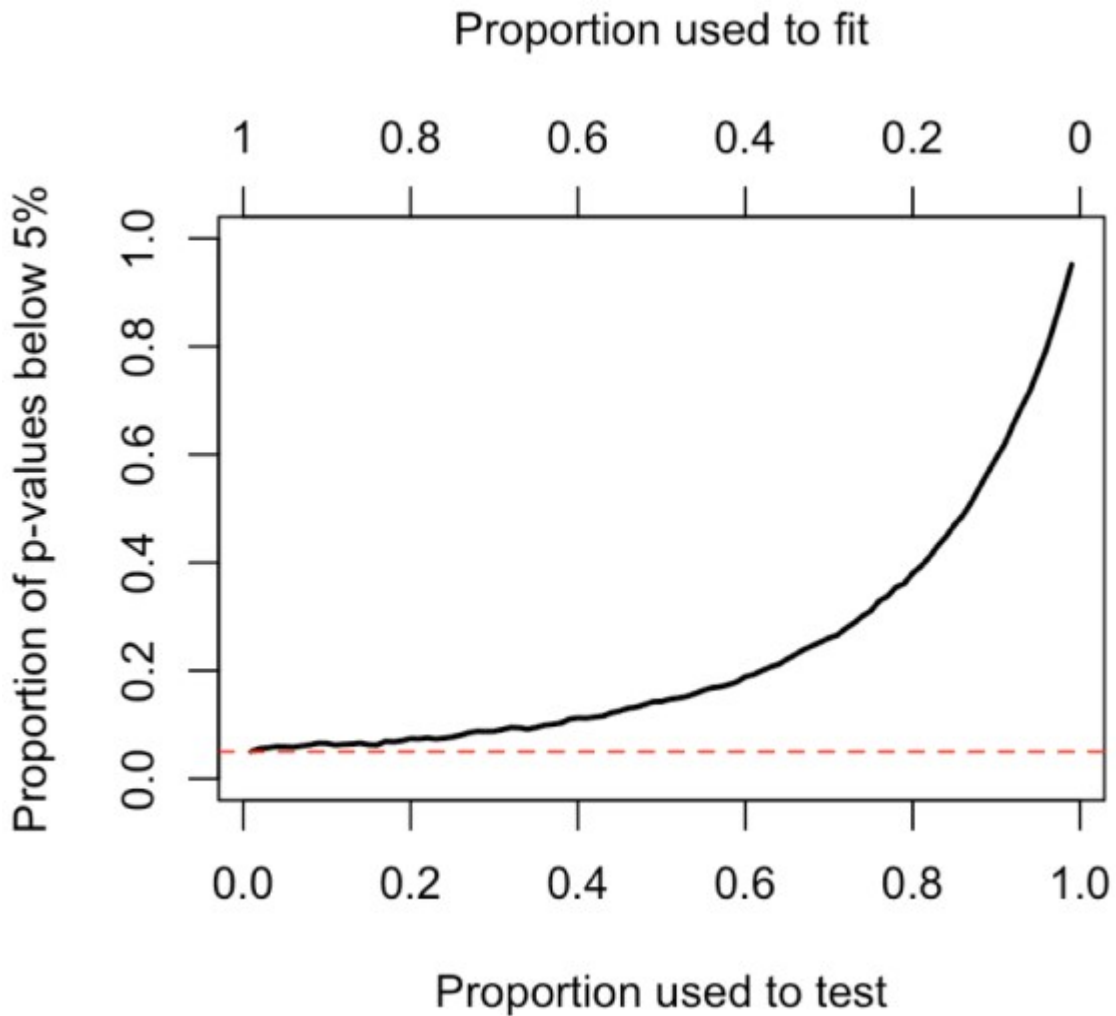
```
mean(p[,3] <= .05)
[1] 0.24168
mean(p[,4] <= .05)
[1] 0.09334
```

We can actually compute that probability as a function of k/n

```
n=300
p = matrix(NA, 1e4, 99)
for(s in 1:1e4){
  X = rnorm(n, 0, 1)
  KS = function(p) ks.test(X[1:(p*n)], "pnorm",
    mean(X[(p*n+1):n]), sd(X[(p*n+1):n]))$p.value
  p[s,] = Vectorize(KS)((1:99)/100)
}
```

The evolution of the probability is the following

```
prob5pc = apply(p, 2, function(x) mean(x <= .05))
plot((1:99)/100, prob5pc)
```



so,

it looks like we can use some sort of hold-out procedure to test for $H_0: X_i \sim \mathcal{N}(\mu, \sigma^2)$, for some μ and σ^2 , using [Kolmogorov–Smirnov test](#) with $\mu = \hat{\mu}$ and $\sigma^2 = \hat{\sigma}^2$ but the proportion of data used to estimate those quantities should be (much) larger than the one used to compute the statistics. Otherwise, we clearly reject too frequently H_0 .